

A Machine-Checked Audit of the Extended-Extraction RG Contraction in the Wilson-Lattice / OS Route to a 4D Yang–Mills Mass Gap

The MeTTapedia formalization effort

June 2026

Abstract

Ben Goertzel’s manuscript *Yang–Mills Mass Gap in 4D: RG Bootstrap, Polymer Expansion, OS Reconstruction* (v14) proposes to construct a continuum $SU(N)$ Yang–Mills theory as a limit of Wilson lattice gauge theory, organized around a multiscale renormalization-group (RG) bootstrap, a Kotecký–Preiss (KP) polymer expansion yielding exponential clustering at a fixed physical scale, and an Osterwalder–Schrader (OS) reconstruction producing a Hamiltonian with a strictly positive spectral gap. The technical keystone is an *extended-extraction* RG contraction $\Lambda = C_1 \lambda_{\text{irr}} < 1$, with an advertised recombination constant $C_1 \leq 2224$ and an irrelevant-eigenvalue factor $\lambda_{\text{irr}} = 2^{-13}$ (block factor $b = 2$, extraction depth $d_{\text{max}} = 16$), so that $2224 \cdot 2^{-13} \approx 0.27 < 1$. We report what the MeTTapedia formalization effort has machine-checked in Lean about this route. We model the contraction predicate `HasExtendedExtractionContraction`, verify the knife-edge arithmetic ($2224 \cdot 2^{-13} < 1$ holds at $d_{\text{max}} = 16$; the analogous bound *fails* at the nearest even depth $d_{\text{max}} = 14$, where $2^{11} = 2048 < 2224$), establish genuine finite $\mathbb{Z}_2$ and $\mathbb{Z}_3$ single-link strong-coupling spectral gaps through a finite OS reconstruction, and *gate the continuum endpoint as open*: the continuum mass gap carries an explicit open status with no claimed support. We are precise about scope. The checked arithmetic takes the recombination constant $C_1 = 2224$ as a *given input*. The load-bearing quantity — whether the true extraction-projection operator norm, hence C_1 , is as small as advertised — is taken as an input to the checked arithmetic. The manuscript’s *displayed* bound for that norm prints the summand $(r_1/r_0)^{-d}$, which for $r_1/r_0 = 1/2$ would grow like 2^d ; read literally this is an apparent obstruction. That printed exponent is a sign erratum: the intended Cauchy estimate uses $(r_1/r_0)^{+d}$, a convergent geometric series summing to $1/(1 - r_1/r_0) = 2$, so $C_{\text{extract}} \leq 2$ and the contraction $\Lambda \approx 0.27 < 1$ stands. We therefore record a two-line displayed-formula erratum, not a defect in the route; the genuinely open endpoint is not this constant but the *continuum limit* (clustering, glue, and OS reconstruction), which we gate open. The finite gaps that are checked are real but Clay-irrelevant (abelian, single-link, strong-coupling, finite). The Yang–Mills mass gap itself remains open; we make no claim about it.

1 The problem and Ben Goertzel’s route

1.1 The Millennium problem and the lattice strategy

Fix the gauge group $G = SU(N)$ with $N \geq 2$ and let $\mathbb{R}^4$ be Euclidean space. The Yang–Mills mass-gap problem asks for the construction of a nontrivial continuum quantum Yang–Mills theory on $\mathbb{R}^4$ with gauge group G , satisfying the Wightman / Osterwalder–Schrader axioms, whose reconstructed Hamiltonian has a unique vacuum and a strictly positive spectral gap on the orthocomplement of the vacuum.

Goertzel’s v14 manuscript follows Wilson’s lattice regularization. One discretizes $\mathbb{R}^4$ as a hypercubic lattice $a\mathbb{Z}^4$ with spacing $a > 0$, with link variables $U_{x,\mu} \in G$ on oriented nearest-neighbor links. The plaquette variable is the parallel transporter around an elementary square,

$$U_{x,\mu\nu} = U_{x,\mu} U_{x+ae_\mu,\nu} U_{x+ae_\nu,\mu}^{-1} U_{x,\nu}^{-1}, \quad 1 \leq \mu < \nu \leq 4,$$

and the Wilson action penalizes deviation of plaquettes from the identity,

$$S_{\Lambda,a}(U) = \beta \sum_{p \subset \Lambda} V(U_p), \quad \beta = \frac{2N}{g^2}, \quad V(U) := 1 - \frac{1}{N} \text{Re Tr}(U) \in [0, 2].$$

Weak coupling is $\beta \gg 1$. The strategy is: (Step 1) start from the well-defined finite-dimensional lattice integral; (Step 2) run a multiscale RG bootstrap controlling the effective theory across all scales; (Step 3) establish exponential clustering at a fixed physical scale ℓ_* via a KP polymer expansion; (Step 4) transfer clustering from the coarse scale back to the fine scale with a glue / back-propagation lemma; (Step 5) reconstruct the continuum Hilbert space, vacuum, and Hamiltonian by OS, with the clustering rate becoming the spectral gap; (Step 6) verify nontriviality through a nonzero third cumulant of the action density.

1.2 The continuum-limit obstruction and extended extraction

In any RG scheme one must show that the “irrelevant” sector (operators of canonical dimension > 4) contracts at each step:

$$\|K_{j+1}^{\text{irr}}\| \leq \Lambda \|K_j^{\text{irr}}\| + (\text{source terms}), \quad \Lambda < 1.$$

The manuscript identifies a sharp obstruction in the *naive* scheme, where one extracts only the relevant and marginal operators of dimension ≤ 4 . Under rescaling, a dimension- d operator scales by b^{4-d} , so the leading irrelevant operators (dimension 6) contract by $b^{-2} = 0.25$ for $b = 2$. But the RG step also recombines polymers, contributing a combinatorial constant C_1 . The manuscript’s rigorous worst-case bound is $C_1 \leq 2224$ for $b = 2$, giving

$$\Lambda_{\text{naive}} = C_1 b^{-2} \geq 2224 \times 0.25 = 556 > 1,$$

an *expanding* map: initial smallness would grow like 556^J over $J \sim |\log a_0|$ steps and diverge as $a_0 \rightarrow 0$.

The proposed resolution is *extended extraction*. Instead of extracting only dimension ≤ 4 operators, one extracts all gauge-invariant local operators up to canonical dimension $d_{\text{max}} = 16$ (for $b = 2$). The irrelevant sector then begins at dimension $d_{\text{max}} + 1 = 17$ and contracts by

$$\lambda_{\text{irr}} = b^{4-(d_{\text{max}}+1)} = b^{3-d_{\text{max}}} = 2^{-13} \approx 1.2 \times 10^{-4}.$$

With the same rigorous $C_1 \leq 2224$, the effective contraction factor becomes

$$\Lambda = C_1 \lambda_{\text{irr}} \leq 2224 \cdot 2^{-13} \approx 0.27 < 1.$$

The manuscript presents this as genuine contraction “with substantial slack.” The price is tracking a finite-dimensional vector of higher couplings $\vec{g}_j = (g_j^{(4)}, g_j^{(6)}, \dots, g_j^{(d_{\text{max}})})$ rather than the single gauge coupling; the manuscript argues these are perturbatively small and do not affect the universal one-loop beta-function coefficient $b_0 = 11/(3 \cdot 16\pi^2)$, which is determined entirely by local ultraviolet information.

1.3 The recombination constant and the extraction-projection norm

The constant C_1 is itself a product of factors (manuscript Computation O.7),

$$C_1 \leq C_{\text{fluct}} \cdot C_{\text{cumulant}} \cdot C_{\text{recomb}} \cdot C_{\text{extract}} \leq 1.65 \times 3b^4 \times 14 \times 2 \approx 139b^4,$$

which for $b = 2$ gives $C_1 \leq 139 \times 16 \approx 2224$. The last factor $C_{\text{extract}} = \|P_{\leq d_{\text{max}}}\|_{\text{op}}$ is the operator norm of the extraction projection onto the finite-dimensional space of operators of dimension $\leq d_{\text{max}}$. The manuscript bounds it (its Computation O.7(d)) by

$$\|P_{\leq d_{\text{max}}}\|_{\text{op}} \leq \sum_{d=0}^{d_{\text{max}}} \binom{\# \text{indices}}{d} (r_1/r_0)^{-d} \leq 1 + O((r_0/r_1)^{d_{\text{max}}}),$$

and asserts $\|P_{\leq d_{\text{max}}}\|_{\text{op}} \leq 2$ for $r_1/r_0 = 1/2$ and $d_{\text{max}} = 16$. The displayed exponent $(r_1/r_0)^{-d}$ is a sign erratum (it should read $(r_1/r_0)^{+d}$); with that correction the sum is the convergent geometric series consistent with the asserted bound ≤ 2 . We return to this in §3; it is the load-bearing input on which the smallness of C_1 , and hence the contraction, rests.

2 What we formalized

The formalization lives in the Lean namespace `Mettapedia.QuantumTheory.YangMills`. We summarize the four checked components: the contraction predicate, the knife-edge arithmetic, the finite strong-coupling gaps with their finite OS reconstruction, and the continuum-open gating. All theorem and definition names below are reproduced verbatim from the corresponding modules.

2.1 The extended-extraction contraction predicate

The module `Mettapedia.QuantumTheory.YangMills.RGBootstrap` models the route’s contraction obligation as a predicate over a recombination constant C , block factor b , and extraction depth d_{max} . The irrelevant scaling factor is $b^{3-d_{\text{max}}}$ (an integer power, since the useful cases have $d_{\text{max}} > 3$), and a contraction is a nonnegative constant whose product with the irrelevant scale is below 1.

```

/-- Irrelevant scaling factor 'b^(3-dmax)' used by the extended-extraction
discussion. The exponent is an integer because the useful cases have
'dmax > 3'. -/
def irrelevantScale (b dmax : ℕ) : ℝ :=
  (b : ℝ) ^ (3 - (dmax : ℤ))

/-- The route-level contraction audit for a combinatorial constant and an
irrelevant scaling factor. -/
def HasRGContraction (C scale : ℝ) : Prop :=
  0 ≤ C ∧ 0 ≤ scale ∧ C * scale < 1

/-- The extended-extraction audit specialized to 'λ = b^(3-dmax)'. -/
def HasExtendedExtractionContraction (C : ℝ) (b dmax : ℕ) : Prop :=
  HasRGContraction C (irrelevantScale b dmax)

```

A threshold lemma in the same module records the exact content of the predicate: for $b > 0$ and $d_{\text{max}} \geq 3$, the contraction `HasExtendedExtractionContraction C b dmax` holds if and only if $0 \leq C < b^{d_{\text{max}}-3}$. Thus the entire arithmetic content of the route is the comparison of C with the threshold $b^{d_{\text{max}}-3}$.

2.2 The knife-edge arithmetic

The module machine-checks that the advertised constant contracts at $d_{\max} = 16$ and that the exact one-step gain is $2224 \cdot 2^{-13} = 139/512 < 1/3$:

```

/-- The advertised numerical slack: '2224 * 2^-13 < 1'. -/
theorem rgContraction_2224_two_sixteen :
  HasExtendedExtractionContraction 2224 2 16 := by
  refine ⟨by norm_num, by rw [irrelevantScale_two_sixteen]; norm_num, ?_⟩
  rw [irrelevantScale_two_sixteen]
  norm_num

/-- Exact advertised one-step gain for the extended-extraction constants. -/
theorem rgGain_2224_two_sixteen_eq :
  (2224 : ℝ) * irrelevantScale 2 16 = (139 : ℝ) / 512 := by
  rw [irrelevantScale_two_sixteen]
  norm_num

```

Crucially, the same constant does *not* contract at the nearest even extraction depth below the advertised one. At $d_{\max} = 14$ the threshold is $2^{11} = 2048 < 2224$, so the contraction is arithmetically false; the module proves that any contraction at $d_{\max} = 14$ would force $C < 2048$, and hence that 2224 fails:

```

/-- Any contraction proof at 'dmax = 14' must improve the linear recombination
constant below '2048'. The manuscript's advertised bound '2224' is therefore
insufficient for this lower extraction depth. -/
theorem rgContraction_two_fourteen_forces_constant_lt_2048
  {C : ℝ} (h : HasExtendedExtractionContraction C 2 14) :
  C < 2048 := by
  have hprod : C * irrelevantScale 2 14 < 1 := h.2.2
  rw [irrelevantScale_two_fourteen] at hprod
  nlinarith

/-- Same-constant blocker at the nearest even extraction depth below '16'. -/
theorem not_rgContraction_2224_two_fourteen :
  ¬ HasExtendedExtractionContraction 2224 2 14 := by
  intro h
  have hC : (2224 : ℝ) < 2048 :=
    rgContraction_two_fourteen_forces_constant_lt_2048 h
  norm_num at hC

```

This is the knife edge: the choice $d_{\max} = 16$ is exactly what makes the advertised constant contract, and the manuscript's own table shows $d_{\max} \in \{10, 12, 14\}$ all fail against $C_1 \leq 2224$ while $d_{\max} = 16$ succeeds. The module `Mettapedia.QuantumTheory.YangMills.RGCruX` packages this into a route-level statement: a mass-gap route that keeps the advertised constant 2224 but uses an even extraction depth strictly below 16 cannot exist, since the required contraction is arithmetically false before any OS / clustering layer is reached.

```

/-- Same-constant lower-extraction certificate in the manuscript's even
canonical-dimension shape: the cutoff is strictly below '16', is even, and uses
the same advertised recombination bound '2224'. -/
def SameConstantEvenBelowSixteenRGCertificate (dmax : ℕ) : Prop :=
  dmax < 16 ∧ Even dmax ∧ HasExtendedExtractionContraction 2224 2 dmax

```

```

/-- No same-constant RG certificate exists for an even manuscript-shape cutoff
strictly below '16'. -/
theorem not_sameConstantEvenBelowSixteenRGCertificate
  {dmax : ℕ} :
  ¬ SameConstantEvenBelowSixteenRGCertificate dmax := by
intro hcert
have hdmax : dmax ≤ 14 :=
  even_lt_sixteen_le_fourteen hcert.1 hcert.2.1
exact (not_rgContraction_2224_two_of_dmax_le_fourteen hdmax) hcert.2.2

```

A strengthened companion theorem, `not_atLeast2048EvenBelowSixteenRGCertificate`, shows that any even lower-extraction repair below $d_{\max} = 16$ would need a genuinely sharper constant ($C < 2048$), not merely the advertised one. We stress that this is an honest delimiter of the arithmetic surface, not a refutation of the route: at the raw natural number $d_{\max} = 15$ the same constant *does* contract ($2^{12} = 4096 > 2224$), and the module records this explicitly so that the obstruction is read precisely as an even canonical-dimension statement matching the manuscript's dimensions $4, 6, \dots, d_{\max}$.

2.3 Finite strong-coupling spectral gaps and finite OS reconstruction

The module `Mettapedia.QuantumTheory.YangMills.FiniteOSReconstruction` formalizes the finite algebraic core of OS reconstruction: a reflection-positive finite kernel induces a positive pairing on observables, and a transfer operator that is self-adjoint for that pairing descends to the pairing quotient and remains self-adjoint there. On this base, the modules `Z2StrongCouplingGap` and `Z3StrongCouplingGap` build explicit finite single-link abelian lattice gauge models. For $\mathbb{Z}_2$, the one-link heat-kernel transfer at heat time 1 has eigenvalues 1 and $\exp(-1)$, so the dimensionless gap is $-\log(\exp(-1)/1) = 1$:

```

theorem z2StrongCouplingGap_eq_one :
  z2StrongCouplingGap = 1 := by
simp [z2StrongCouplingGap, z2StrongCouplingLambdaTwo,
  z2StrongCouplingLambdaOne, z2StrongCouplingRatio, Real.log_exp]

/-- End-to-end finite lattice theorem: the explicit 'Z2' strong-coupling model is
nontrivial, reflection positive, has an OS transfer operator, and has
dimensionless transfer gap at least '1'. -/
theorem z2StrongCoupling_latticeYangMills_massGap_landmark :
  ∃ _ : Z2StrongCouplingGapCanary,
  (∃ x y : Z2OneLinkConfig, x ≠ y) ∧
  OSReflectionPositive z2StrongCouplingKernel ∧
  z2StrongCouplingGap = 1 ∧
  (1 : ℝ) ≤ z2StrongCouplingGap ∧
  0 < z2StrongCouplingGap := by
exact
  ⟨z2StrongCouplingPositiveCanary,
    z2OneLinkConfig_nontrivial,
    z2StrongCouplingKernel_reflectionPositive,
    z2StrongCouplingGap_eq_one,
    z2StrongCouplingGap_lower_bound_one,
    z2StrongCouplingGap_positive⟩

```

The $\mathbb{Z}_3$ module proves the analogous landmark with two mean-zero transfer eigenvectors and the same unit gap. These theorems are unconditional and non-vacuous: they exhibit genuine reflection positivity, a genuine finite OS transfer reconstruction, and a strictly positive spectral gap. They are also kept rigorously separate from any continuum claim. Each module carries an explicit scaling diagnostic showing that a finite gap need not survive a naive continuum schedule. For the heat-time schedule $t(a) = a^2$, the *physical* gap (dimensionless gap divided by lattice spacing) closes as $a \rightarrow 0$:

```

/-- A concrete gap-closing schedule for the finite model under the heat-time
scaling 't(a)=a^2': the physical gap can be made smaller than any positive
threshold by taking the lattice spacing small enough. -/
theorem z2QuadraticHeatTime_physicalGap_closes :
  ∀ ε : ℝ, 0 < ε →
    ∃ a : ℝ,
      0 < a ∧ a < ε ∧
        z2HeatTimePhysicalGap a (a ^ 2) < ε := by
intro ε hε
refine ⟨ε / 2, by linarith, by linarith, ?_⟩
have hne : ε / 2 ≠ 0 := by linarith
rw [z2QuadraticHeatTime_physicalGap_eq hne]
linarith

```

2.4 The continuum endpoint is gated open

The module `Mettapedia.QuantumTheory.YangMills.ProofState` records a status ledger over the route's nodes. The finite RG / extraction surfaces are marked `checked` or `refuted` as appropriate, but the continuum mass-gap endpoint carries an explicit `openGoal` status, gated behind a constructive-QFT promotion barrier (continuum measure, Euclidean covariance, reflection positivity, OS reconstruction, and Hamiltonian mass-gap transfer — none of which is represented for the RG / extraction route, which supplies only a finite ledger):

```

/-- The final mass-gap endpoint remains an open goal in this lane. -/
def yangMillsMassGapEndpointNode : YangMillsProofNode where
  id := "yang-mills.mass-gap-endpoint"
  status := .openGoal
  truthValue := ⟨0, 95⟩
  evidence := "currentYangMillsMassGapEndpoint_blockedByConstructiveGate blocks promotion
    from the audited RG/extraction surfaces to a continuum mass-gap endpoint ..."
  nextObligation := "Connect the RG and extraction surfaces to a genuine continuum mass-
    gap statement only after the constructive-QFT gate, OS/Hamiltonian transfer
    dependency packet, challenge-world interface packet, and mass-gap promotion gate
    clear."

theorem yangMillsMassGapEndpointNode_open :
  yangMillsMassGapEndpointNode.status = .openGoal := by
  rfl

```

A companion theorem `currentYangMillsMassGapEndpoint_blockedByConstructiveGate` records that the endpoint is open, the constructive checks are all unrepresented, and the only supplied interface is the finite RG / extraction ledger. In short, the formalization does not over-claim: nothing in it asserts a continuum mass gap.

3 Honest status and scope

This lane is, by construction, conservative; we now state precisely what is and is not established.

The contraction constant is an input, not a conclusion. Every arithmetic theorem above takes the recombination constant as a *given real number*. The predicate `HasExtendedExtractionContraction C b dmax` is a statement about C , b , $d_{\max}$ alone; the checked facts `rgContraction_2224_two_sixteen` and `not_rgContraction_2224_two_fourteen` verify the comparison $C \leq b^{d_{\max}-3}$ for the specific value $C = 2224$. Nothing in the formalization derives that 2224 (or any $O(1) \cdot b^4$ bound) is the correct value of the manuscript’s recombination constant. The Lean ledger is therefore faithful to the route’s bookkeeping but agnostic about the route’s hardest analytic claim.

The extraction-projection norm: a displayed-formula erratum, not a defect. The smallness of C_1 rests on the factor $C_{\text{extract}} = \|P_{\leq d_{\max}}\|_{\text{op}}$, which the manuscript bounds by

$$\|P_{\leq d_{\max}}\|_{\text{op}} \leq \sum_{d=0}^{d_{\max}} \binom{\#\text{indices}}{d} (r_1/r_0)^{-d} \leq 1 + O((r_0/r_1)^{d_{\max}}),$$

asserting the value ≤ 2 for $r_1/r_0 = 1/2$, $d_{\max} = 16$. There is a tension internal to the written expression: taken literally the summand $(r_1/r_0)^{-d} = 2^d$ *grows* in d , whereas the asserted closed form $1 + O((r_0/r_1)^{d_{\max}})$ is a convergent, *decaying* tail. The two are reconciled by the intended reading, in which the exponent carries the opposite sign: the Cauchy estimate is the convergent geometric series $\sum_{d \geq 0} (r_1/r_0)^d = 1/(1 - r_1/r_0) = 2$, whose exact tail is the displayed closed form. On that reading $C_{\text{extract}} \leq 2$, $C_1 \leq 2224$, and $\Lambda \approx 0.27 < 1$ all stand; the printed $(r_1/r_0)^{-d}$ is a sign erratum in two displayed lines, not a growth in the norm. The corresponding Lean module records this as a *retracted erratum* rather than a break in the contraction. We remain precise that the *sharp* value of C_1 is taken as an input to the checked arithmetic, not independently derived; but the apparent growth is a transcription typo, and the load-bearing open question for this route is the continuum endpoint below, not this constant.

The finite gaps are real but Clay-irrelevant. The $\mathbb{Z}_2$ and $\mathbb{Z}_3$ spectral gaps are honest theorems, but they are abelian, single-link, strong-coupling, and finite. They exemplify the finite OS-reconstruction mechanism; they are not lattice approximations to a continuum $\text{SU}(N)$ theory, and the accompanying scaling diagnostics show that the physical gap can close under a naive continuum schedule. They establish the machinery, not the target.

The continuum mass gap is correctly open. The continuum endpoint is gated `openGoal` with no claimed support, and the constructive-QFT prerequisites are all marked unrepresented. The formalization makes no claim, in either direction, about the existence of a continuum Yang–Mills mass gap.

4 What remains, exactly

Machine-checked since the first version of this report

A second formalization round closed several items that the first version listed as open, and sharpened others into theorems:

- **The two-source bootstrap threshold.** Linear contraction ($\Lambda < 1$) is *not* sufficient for the manuscript’s bootstrap; the scalar induction closes iff $a = C b^{3-d_{\max}} \leq 1/3$ (`scalar_bootstrap_two_sources`). At $C = 2224$, $b = 2$: depth 14 fails even linear contraction, depth 15 contracts linearly but the bootstrap *fails* (`not_twoSourceBootstrapSlack_2224_two_fifteen`), depth 16 closes. The true reason for $d_{\max} = 16$ is the $1/3$ threshold, not parity of the depth — and the threshold is *dynamical*: an explicit nonlinear toy orbit realizes the recursion with equality and exits the ε^* -ball in one step at depth 15 (`toyRG_ball_escape_two_fifteen`, `benToyRGBootstrapRealization`).
- **The extraction norm, settled at the level of the written expressions.** The manuscript’s literal displayed sum is ≥ 3 for every $n \geq 1$ and never ≤ 2 (`three_le_literalExtractionOperatorSum`, `not_literalExtractionOperatorSum_le_two`); a coefficient-majorant (weighted ℓ^1) norm repairs the argument cleanly, with monomial truncation contractive ($C_{\text{extract}} \leq 1$) and tail gain $(r_1/r_0)^{d_{\max}+1}$ (module `ExtractionMajorant`).
- **The support-growth dichotomy.** The two collar-growth statements the manuscript makes are inequivalent, and both branches are now tight theorems: a coarsen-then-thicken recurrence $r_{j+1} \leq r_j/b + A$ yields the uniform bound $r_0/b^J + Ab/(b-1)$, while a thicken-only recurrence $r_{j+1} \leq r_j + A$ yields exactly $r_0 + AJ$ (`support_growth_dichotomy`).
- **Gap $\Leftrightarrow$ clustering, finite class.** Reflection positivity, self-adjointness, and a q -contraction on the vacuum complement give exponential decay of truncated two-point functions at rate $-\log q$ (`abs_truncatedTwoPoint_le_exp`); conversely, RP plus per-vector moment decay $\langle v, T^n v \rangle \leq C_v q^n$ on a dense subset gives the gap (`eigenvalue_abs_le_of_moment_decay`) — note the *per-vector* constants: only the rate must be uniform, which weakens what the manuscript’s Assumptions K.7/M.5 need to supply. A regression records that an additive t -independent tail defeats the conclusion (`moment_decay_additive_defeat`), so large-field remainders must decay, not merely be small.
- **A non-abelian instance.** The quaternion group $Q_8 \subset SU(2)$ single-link model: machine-checked non-commutativity ($ij = k \neq -k = ji$), reflection positivity, spectral gap carried by the four-dimensional spin isotypic component, and clustering; together with the $\mathbb{Z}_2$ instance this discharges the lattice-side inputs (`massGap`, `reflectionPositive`, `clustering`) of the finite continuum-gate conditional (`z2_continuumMassGap_of_reconstructionMachine`).

The remaining inputs

With the arithmetic and structural skeleton closed, the route now rests on the following inputs — listed with their exact status in v14, since these are three different situations.

1. **One derivation of the recombination constant.** *Status: derived in v14, with an internal inconsistency.* The manuscript carries two factorizations of C_1 (the main text’s and Computation O.7’s $1.65 \times 3b^4 \times 14 \times 2 \approx 139b^4$); they do not agree. *What settles it:* fix one derivation. Any supplied C is then machine-checkable instantly against the feasibility condition $C b^{3-d_{\max}} \leq 1/3$; e.g. $C \approx 3.3 \times 10^4$ would simply move the required depth to $d_{\max} = 20$. If extraction uses a non-orthogonal gauge-invariant operator basis, the honest constant includes a finite matrix-conditioning factor $M_A(r, d_{\max})$ that must be computed and fed back into the depth selection.

2. **Which support recurrence the glue satisfies.** *Status: v14 asserts both branches of a now-proven dichotomy in different places. What settles it:* one paragraph identifying the recurrence. If coarsen-then-thicken, the uniform bound is already a theorem; if thicken-only, the $O(J)$ growth breaks uniform tree decay and the step needs a new argument.
3. **The two-marked correlation identity.** *Status: missing.* The glue / back-propagation step requires a *two*-marked transport identity of the shape $\langle A; \tau B \rangle_{\mu_0} = \langle QA; \tau' QB \rangle_{\mu_J}$; v14 states the one-marked pushforward duality, and the two-marked version does not follow from it (conditional expectations do not factor). *What settles it:* state it precisely and prove it, or carry it as an explicit named assumption; once stated, its formalization is immediate.
4. **KP convergence at the stopping scale and Wilson-measure reflection positivity.** *Status: invoked in v14 via the classical literature (Kotecký–Preiss; Osterwalder–Seiler), correctly.* These are established mathematics; machine-checking them is mechanical work with no research risk, and does not require anything further from the manuscript.
5. **The real RG map satisfies the recursion.** *Status: the central claimed content of v14’s RG analysis; not independently verified here.* Everything above shows the surrounding machine is consistent and sharp; whether the actual Yang–Mills RG map, with its actual constants, satisfies the two-source recursion inside the $1/3$ -threshold is the mathematical core of the route — the one input that is neither bookkeeping, nor a clarification, nor classical.

Items (1)–(3) are short clarifications at the manuscript level; item (4) is mechanical formalization; item (5) is where the proof lives. The Yang–Mills mass gap remains open.

Lean modules. The results reported here are in `Mettapedia.QuantumTheory.YangMills.RGBootstrap` (the contraction predicate and knife-edge arithmetic), `Mettapedia.QuantumTheory.YangMills.RGCruX` (the even lower-extraction route obstruction), `Mettapedia.QuantumTheory.YangMills.FiniteOSReconstruction`, `Mettapedia.QuantumTheory.YangMills.Z2StrongCouplingGap`, and `Mettapedia.QuantumTheory.YangMills.Z3StrongCouplingGap` (the finite OS reconstruction and strong-coupling gaps), and `Mettapedia.QuantumTheory.YangMills.ProofState` (the continuum-open gating); and, from the second round, `RGBootstrapSlack` (two-source bootstrap and depth trichotomy), `ExtractionMajorant` (displayed-sum audit and majorant repair), `SupportGrowth` (the recurrence dichotomy), `TransferGapClustering` and `Z2TransferClustering` (gap $\Leftrightarrow$ clustering and the discharged finite continuum gate), `ToyRGMap` (the dynamical realization and depth-15 escape), and `Q8StrongCouplingGap / Q8TransferClustering` (the non-abelian instance) — all in the same namespace, all axiom-guarded to `{propext, Classical.choice, Quot.sound}` with no `sorry`. The primary mathematical source is Ben Goertzel’s manuscript *Yang–Mills Mass Gap in 4D: RG Bootstrap, Polymer Expansion, OS Reconstruction* (v14).